



SAIR

Spatial AI & Robotics Lab

CSE 473/573

L19: FACE DETECTION

Chen Wang


Spatial AI & Robotics Lab

Department of Computer Science and Engineering



University at Buffalo The State University of New York

Why is detection hard?



We want to do this
for ~ 1000 objects

1,000,000 images/day



Why is detection hard?


If we have 1000 categories (detectors), and each detector produces 1 false every 10 images, we will have 100 false alarms per image... pretty much garbage...



Local information is helpful



Is local information enough?



Is local information even enough?

We know there is a keyboard present in this scene even if we cannot see it clearly.







We know there is no keyboard present in this scene



... even if there is one indeed.

The multiple personalities of a blob



The multiple personalities of a blob

A 13 C

12
13
14

12
A 13 C
14

The context challenge

- What are the hidden objects?



The context challenge

- What are the hidden objects?





Object detection vs Scene Recognition


- Objects (even if deformable and articulated) probably have more **consistent shapes** than scenes.
- Scenes can be defined by distribution of “stuff” – materials and surfaces with **arbitrary shape**.
- Objects are “**things**” that **own their boundaries**
- Bag of words (BoW) were less popular for object detection because they **throw away shape** info.

Object Category Detection

- Focus on object search: “Where is it?”
- Build templates that quickly differentiate object patch from background patch



Dog Model



**Object or
Non-Object?**

Challenges in modeling the object class



Illumination



Object pose




Clutter



Occlusions



Intra-class
appearance




Viewpoint

Challenges in modeling the non-object class

True
Detections




Bad
Localization



Confused with
Similar Object




Misc. Background



Confused with
Dissimilar Objects



General Process of Object Recognition



What are the object parameters?


Specifying an object model

1. Statistical Template in Bounding Box

- Object is somewhere (x,y,w,h) in image
- Features defined wrt bounding box coordinates



Image





Template Visualization
(HoG features)

Specifying an object model

2. Articulated parts model


- Object is configuration of parts
- Each part is detectable









Specifying an object model

3. Hybrid template/parts model

Detections



Template Visualization
HoG features



root filters
coarse resolution


part filters
finer resolution

deformation
models


Specifying an object model

4. 3D-ish model

- Object is collection of 3D planar patches under affine transformation



General Process of Object Recognition



Generating hypotheses

1. Sliding window

- Test patch at each **location** and **scale**



Generating hypotheses

1. Sliding window

- Test patch at each **location** and **scale** (pyramids)



Note – Template did not change size

Each window is separately classified




Generating hypotheses

2. Voting from patches/keypoints



Interest Points




Matched Codebook Entries



Probabilistic Voting





3D Voting Space (continuous)




Generating hypotheses

3. Region-based proposal




General Process of Object Recognition




Mainly-gradient based features, usually based on summary representation, many classifiers.

HOG + SVM for Object Detection




<https://youtu.be/sDByl84n5mY>

General Process of Object Recognition





Rescore each proposed object based on whole set


Resolving detection scores

$$IOU = \frac{\text{area of overlap}}{\text{area of union}}$$



1. Non-max suppression



Resolving detection scores

$$IOU = \frac{\text{area of overlap}}{\text{area of union}}$$



1. Non-max suppression



“Overlap” score is below some threshold

Resolving Detection Scores



2. Context/reasoning



(g) Car Detections: Local



(h) Ped Detections: Local




Non-Max Suppression

Confidence Scores



0 → 100



Set a high confidence score threshold of 95%



Face detection and recognition



Detection



Recognition



"Sally"

Consumer application:



Consumer application:

- Can be trained to recognize pets!



Challenges of face detection

- Sliding window detector must evaluate tens of thousands of location/scale combinations
- Faces are rare: 0–10 per image
 - For efficiency, we should try to spend as little time as possible on the non-face windows
 - A megapixel image has $\sim 10^6$ pixels and a comparable number of candidate face locations
 - To avoid having a false positive in every image, our false positive rate has to be less than $1/10^6$



The Viola/Jones Face Detector

- Viola-Jones Detection Framework
 - Sliding Window Face Detection
- A simple but effective approach to real-time object detection
- Training is slow, but detection is very fast
- Key ideas
 - *Integral images* for fast feature evaluation
 - *Boosting* for feature selection
 - *Attentional cascade* for fast rejection of non-face patches

Image Features

- Simple Features that measures the difference in intensity


“Rectangle filters”




$$Value = \sum (pixels\ in\ white) - \sum (pixels\ in\ black)$$

Example

Source




Result




Integral Image (Recap)

- A transformed image where every pixel is the sum of all pixels **above** and to the **left** of original image.




Integral images (Recap)

- What's the sum of pixels in the Rectangle ABCD?




Computing sum within a rectangle (Recap)

- Let A,B,C,D be the values of the integral image at the corners of a rectangle
- Then the sum of original image values within the rectangle can be computed as:
$$\text{sum} = D - B - C + A$$
- Only 3 additions are required for any size of rectangle!





Integral Image Example (Recap)



Computing a rectangle feature



- Verify this:
 - Value = $\text{sum}(\text{scale factors} * \text{area of the regions})$.



Exercise
10 minutes

How do we get this?

- White Area – Black Area
- Sum of Entire Block = $A - B - C + D$
- White Block = $A - F - C + E$
- Black Block = $F - E - B + D$
- White – Black =
 - $A - 2F - C + 2E - D + B$




Key ideas

- *Integral images* for **fast feature evaluation**
- ***Boosting*** for feature selection
- *Attentional cascade* for fast rejection of non-face windows

Feature selection

- For a 24x24 detection region, the number of possible rectangle features is ~160,000!



Feature selection

- At test time, it is impractical to evaluate the entire feature set
- Can we create a good classifier using just a small subset of all possible features?
- How to select such a subset?

Boosting

- *Boosting* is a learning scheme that combines *weak learners* into a more accurate *ensemble classifier*
- Weak learners based on rectangle filters:

value of rectangle feature

↓

$$h_t(x) = \begin{cases} 1 & \text{if } p_t f_t(x) > p_t \theta_t \\ 0 & \text{otherwise} \end{cases}$$

Annotations:

- ↑ window (points to $h_t(x)$)
- ↑ parity (points to p_t)
- ↑ threshold (points to θ_t)

- Ensemble classification function:

$$C(x) = \begin{cases} 1 & \text{if } \sum_{t=1}^T \alpha_t h_t(x) > \frac{1}{2} \sum_{t=1}^T \alpha_t \\ 0 & \text{otherwise} \end{cases}$$


Annotations:

- ← learned weights (points to α_t)

A parity indicates the direction of the inequality sign


Training procedure

- Initially, weight each training example equally
- In each boosting round:
 - Find the weak learner with lowest *weighted* training error
 - Raise the weights of training examples misclassified by current weak learner
- Compute final classifier as linear combination of all weak learners
 - Weight of each learner is proportional to its accuracy
 - Exact formulas for re-weighting and combining weak learners depend on the particular boosting scheme, i.e. Adaptive Boost (AdaBoost).




Boosting intuition

**Weak
Classifier 1**




Boosting illustration

**Weights
Increased**




Boosting illustration




Boosting illustration

**Weights
Increased**




Boosting illustration




Boosting illustration

**Final classifier is
a combination of weak
classifiers**



Training procedure

- Initially, weight each training example equally
- In each boosting round:
 - Find the weak learner with lowest *weighted* training error
 - Raise the weights of training examples misclassified by current weak learner
- Compute final classifier as linear combination of all weak learners
 - Weight of each learner is proportional to its accuracy
 - Exact formulas for re-weighting and combining weak learners depend on the particular boosting scheme, i.e. Adaptive Boost (AdaBoost)



Adaboost for face detection

- Given example images $(x_1, y_1), \dots, (x_n, y_n)$ where $y_i = 0, 1$ for negative and positive examples respectively.
- Initialize weights $w_{1,i} = \frac{1}{2m}, \frac{1}{2l}$ for $y_i = 0, 1$ respectively, where m and l are the number of negatives and positives respectively.
- For $t = 1, \dots, T$:

1. Normalize the weights,

$$w_{t,i} \leftarrow \frac{w_{t,i}}{\sum_{j=1}^n w_{t,j}}$$

so that w_t is a probability distribution.

2. For each feature, j , train a classifier h_j which is restricted to using a single feature. The error is evaluated with respect to w_t , $\epsilon_j = \sum_i w_i |h_j(x_i) - y_i|$.

3. Choose the classifier, h_t , with the lowest error ϵ_t .
4. Update the weights:

$$w_{t+1,i} = w_{t,i} \beta_t^{1-e_i}$$

where $e_i = 0$ if example x_i is classified correctly, $e_i = 1$ otherwise, and $\beta_t = \frac{\epsilon_t}{1-\epsilon_t}$.


- The final strong classifier is:

$$h(x) = \begin{cases} 1 & \sum_{t=1}^T \alpha_t h_t(x) \geq \frac{1}{2} \sum_{t=1}^T \alpha_t \\ 0 & \text{otherwise} \end{cases}$$

where $\alpha_t = \log \frac{1}{\beta_t}$

Boosting for face detection

- First two features selected by boosting:
- This feature combination can yield 100% recall and 50% false positive rate



Boosting vs. SVM


- Advantages of boosting
 - Integrates classifier training with feature selection
 - Complexity of training is linear instead of quadratic in the number of training examples
 - Flexibility in the choice of weak learners, boosting scheme
 - Testing is fast
- Disadvantages
 - Needs many training examples
 - Training is slow
 - Many-class problems doesn't work as well as SVM

Key ideas

- *Integral images* for **fast feature evaluation**
- *Boosting* for feature selection
- *Attentional **cascade*** for fast rejection of non-face windows


Boosting for face detection

- A 200-feature classifier can yield 95% detection rate and a false positive rate of 1/14084




Attentional cascade

- Start with simple classifiers which reject many negatives while detecting almost all positives.
- Positive response from the first classifier triggers the evaluation of a second (more complex) classifier, etc.
- A negative outcome at any point leads to the immediate rejection of the sub-window





Attentional cascade

- Chain classifiers that are progressively more complex and have lower false positive rates




ROC




Attentional cascade

- The detection rate and the false positive rate of the cascade are found by multiplying the respective rates of the individual stages
- A detection rate of 0.9 and a false positive rate 10^{-6} can be achieved by a 10-stage cascade if each stage has a detection rate of 0.99 ($0.99^{10} \approx 0.9$) and a false positive rate of about 0.30 ($0.3^{10} \approx 6 \times 10^{-6}$)



Training the cascade

- Set target detection and false positive rates for each stage
- Keep adding features to current stage until target rates met
 - Lower AdaBoost threshold to maximize detection
 - opposed to minimizing total classification error
 - Test on a *validation set*
- If the overall false positive rate is not low enough
 - Add another stage
- Use false positives from current stage as the negative training examples for the next stage



The implemented system





- Training Data
 - 5000 faces
 - All frontal, rescaled to 24x24 pixels
 - 300 million non-faces
 - 9500 non-face images
 - Faces are normalized
 - Scale, translation
- Many variations
 - Across individuals
 - Illumination
 - Pose



System performance

- Training time: “weeks” on 466 MHz Sun workstation
- 38 layers, total of 6061 features
- Average of 10 features evaluated per window on test set
- “On a 700 Mhz Pentium III processor, the face detector can process a 384 by 288 pixel image in about .067 seconds”
 - 15 Hz
 - 15 times faster than previous detector of comparable accuracy (Rowley et al., 1998)


Output of Face Detector on Test Images



Other detection tasks




Facial Feature Localization



Profile Detection


Male vs.
female



Profile Detection



Profile Features



Summary: Viola/Jones detector

- Rectangle features
- Integral images for fast computation
- Boosting for feature selection
- Attentional cascade for fast rejection of negative windows